

Deployment 5 Documentation

Before Deployment

Create a VPC called kura-vpc with the IPV4 CIDR of 192.168.0.0/16.

Create 4 subnets. 2 private and 2 public with the following ranges.

192.168.0.0 - 192.168.63.255

192.168.64.0 - 192.168.127.255

192.168.128.0 - 192.168.191.255

192.168.192.0 - 192.168.255.255

<input type="checkbox"/>	Name ▾	Subnet ID ▾	State ▾	VPC ▾	IPv4 CIDR
<input type="checkbox"/>	public01	subnet-033073c2283225dd8	✔ Available	vpc-0923116d7924ad7d0 ku...	192.168.0.0/18
<input type="checkbox"/>	private01	subnet-0119f0ebc4a0b7c31	✔ Available	vpc-0923116d7924ad7d0 ku...	192.168.128.0/18
<input type="checkbox"/>	public02	subnet-0660ccbd79e22236e	✔ Available	vpc-0923116d7924ad7d0 ku...	192.168.64.0/18
<input type="checkbox"/>	private02	subnet-041003a6e672f1133	✔ Available	vpc-0923116d7924ad7d0 ku...	192.168.192.0/18

Create an internet gateway to allow our devices to communicate outside of the network. Once created, attach it to the VPC

VPC > Internet gateways > Attach to VPC (igw-04946cd1f6561bf70)

Attach to VPC (igw-04946cd1f6561bf70) [Info](#)

VPC
Attach an internet gateway to a VPC to enable the VPC to communicate with the internet. Specify the VPC to attach below.

Available VPCs
Attach the internet gateway to this VPC.

▶ AWS Command Line Interface command

[Cancel](#) [Attach internet gateway](#)

Once attached, create a private and public route table

Create route table [Info](#)

A route table specifies how packets are forwarded between the subnets within your VPC, the internet, and your VPN connection.

Route table settings

Name - *optional*

Create a tag with a key of 'Name' and a value that you specify.

VPC

The VPC to use for this route table.

Tags

A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.

Key



Value - *optional*



You can add 49 more tags.

Private route tables route traffic to vpc but not outside the internet.

Attach the Internet gateway to the publicRT routes

Edit routes

Destination

Target



Go into the subnet association of publicRT and attach the two public ips to it.

Explicit subnet associations (0)	
<input type="text" value="Find subnet association"/>	
Subnet ID	IPv4 CIDR
Subnets without explicit associations (2)	
The following subnets have not been explicitly associated with any route tables and are therefore associated with the main route table:	
<input type="text" value="Find subnet association"/>	
Subnet ID	IPv4 CIDR
subnet-033073c2283225dd8 / public01	192.168.0.0/18
subnet-0660ccbd79e22236e / public02	192.168.64.0/18

Go into the subnet association of privateRTand attach the two privateips to it.

Explicit subnet associations (2)	
<input type="text" value="Find subnet association"/>	
Subnet ID	IPv4 CIDR
subnet-0119f0ebc4a0b7c31 / private01	192.168.128.0/18
subnet-041003a6e672f1133 / private02	192.168.192.0/18
Subnets without explicit associations (2)	
The following subnets have not been explicitly associated with any route tables and are therefore associated with the main route table:	
<input type="text" value="Find subnet association"/>	
Subnet ID	IPv4 CIDR
subnet-033073c2283225dd8 / public01	192.168.0.0/18
subnet-0660ccbd79e22236e / public02	192.168.64.0/18

Inside of the VPC subnet, select subnets. And actions enable auto sign ip address

Actions ▲

View details

Create flow log

Modify auto-assign IP settings

Create a public EC2 JumpHost

Network ⓘ	vpc-0923116d7924ad7d0 kura-vpc ▼	Create new VPC
Subnet ⓘ	subnet-033073c2283225dd8 public01 us-east-1a ▼ 16378 IP Addresses available	Create new subnet
Auto-assign Public IP ⓘ	Use subnet setting (Enable) ▼	

Assign a security group: ☒ Create a **new** security group
☐ Select an **existing** security group

Security group name:

Description:

Type ⓘ

SSH ▼

Protocol ⓘ

TCP

Port Range ⓘ

22

Source ⓘ

Anywhere ▾


0.0.0.0/0, ::/0

Launch and select the keypair that you have downloaded

Create a private01 EC2 instance

Network ⓘ

vpc-0923116d7924ad7d0 | kura-vpc ▾

 [Create new VPC](#)

Subnet ⓘ

subnet-0119f0ebc4a0b7c31 | private01 | us-east-1a ▾

[Create new subnet](#)

16379 IP Addresses available

Auto-assign Public IP ⓘ

Use subnet setting (Disable) ▾

Assign a security group: ☒ Create a **new** security group

☐ Select an **existing** security group

Security group name:

only_jumphost

Description:

Only_jumphost should have access only

Type ⓘ

SSH ▾

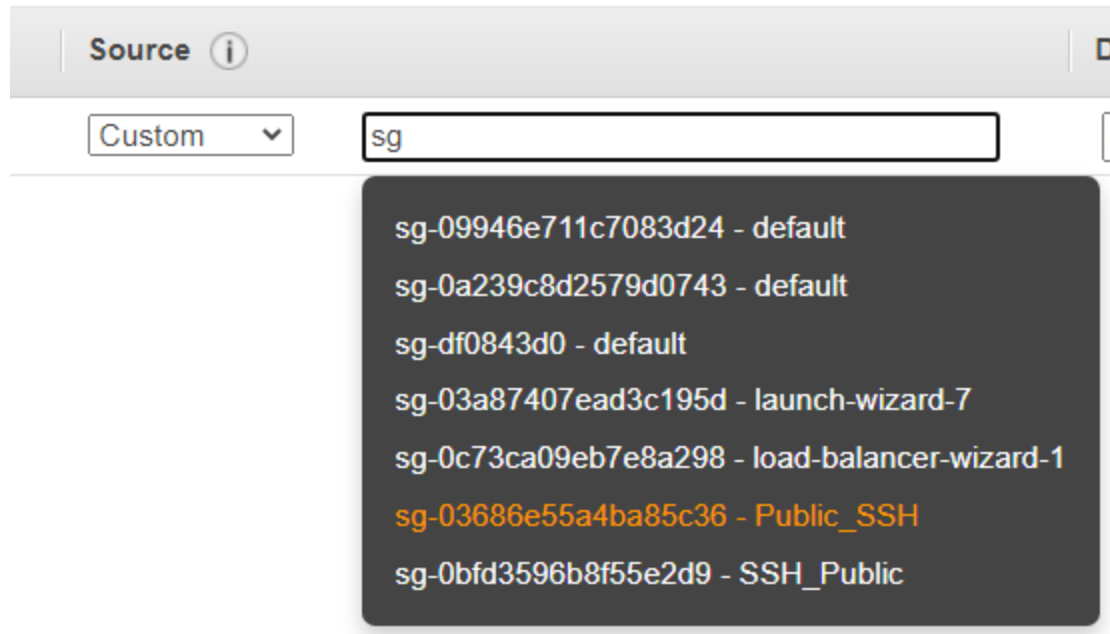
Protocol ⓘ

TCP

Port Range ⓘ

22

Select the security group you created for JumpHost (public ec2) which is Public_SSH



Launch and select the keypair that you have downloaded

Connect to the Public SSH named JumpHost
`ssh -i .key.pem ec2-user@Public IPv4 address`

Run `ping 8.8.8.8`
You should get some requests.

`nano linux.pem`

Now get the information from your pem keypair that you had downaloded and paste it in
linux.pem on the ec2 instance

Save it. Control + O to save and Control + X to exit

`chmod 400 linux.pem`

The Private IPv4 addresses is from the private01 ec2 instance.

`ssh -i linux.pem ec2-user@Private IPv4 addresses`
`ping 8.8.8.8` and you should have 100% packet loss

```
--- 8.8.8.8 ping statistics ---  
10 packets transmitted, 0 received, 100% packet loss, time 9209ms
```

NAT gateway, allows us to restrict inbound but allow us to do outbound.

Go into aws VPC

Create a NAT Gateway on the left side

NAT gateway settings

Name - optional
Create a tag with a key of 'Name' and a value that you specify.

The name can be up to 256 characters long.

Subnet
Select a subnet in which to create the NAT gateway.

subnet-033073c2283225dd8 (public01) ▼

Connectivity type
Select a connectivity type for the NAT gateway.

☒ Public
☐ Private

Elastic IP allocation ID [Info](#)
Assign an Elastic IP address to the NAT gateway.

eipalloc-0674755f1ad7583f2 ▼

Allocate Elastic IP

Go to Routing table

Select the private-RT

Edit the Routes

Add Routes

X

Now go back into terminal and inside the private ec2 that you SSH into, run ping 8.8.8.8

You should see some responses.

```
[ec2-user@ip-192-168-189-227 ~]$ ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data:
64 bytes from 8.8.8.8: icmp_seq=1 ttl=51 time=1.84 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=51 time=0.997 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=51 time=1.04 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=51 time=1.02 ms
64 bytes from 8.8.8.8: icmp_seq=5 ttl=51 time=1.03 ms
64 bytes from 8.8.8.8: icmp_seq=6 ttl=51 time=0.973 ms
```

DEPLOYMENT 5

Step 1

```
sudo amazon-linux-extras install java-openjdk11

sudo amazon-linux-extras install epel

sudo wget -O /etc/yum.repos.d/jenkins.repo \
https://pkg.jenkins.io/redhat-stable/jenkins.repo

sudo rpm --import https://pkg.jenkins.io/redhat-stable/jenkins.io.key

sudo yum upgrade

sudo yum install epel-release java-11-openjdk-devel

sudo yum install jenkins

sudo systemctl start jenkins
```

Step 2

Select Target Group inside of AWS EC2

▼ Load Balancing

Load Balancers

Target Groups New

Then create a target group.

Create target group

Then select instances

Choose a target type

☒ Instances

- Supports load balancing to instances within a specific VPC.

Create a target group name

Target group name

Jenkins

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Select protocol HTTP and put port 8080

Protocol

Port

HTTP



:

8080

Select your VPC

VPC

Select the VPC with the instances that you want to include in the target group.

kura-vpc

vpc-0923116d7924ad7d0

IPv4: 192.168.0.0/16



Select HTTP1

Protocol version

☒ HTTP1

Send requests to targets using HTTP/1.1. Supported when the request protocol is HTTP/1.1 or HTTP/2.

☐ HTTP2

Send requests to targets using HTTP/2. Supported when the request protocol is HTTP/2 or gRPC, but gRPC-specific features are not available.

☐ gRPC

Send requests to targets using gRPC. Supported when the request protocol is gRPC.

Select HTTP and enter /login

Health checks

The associated load balancer periodically sends requests, per the settings below, to the registered targets to test their status.

Health check protocol

HTTP ▼

Health check path

Use the default path of "/" to ping the root, or specify a custom path if preferred.

/login

Up to 1024 characters allowed.

Select Advanced Health check settings and select override and enter 8080

▼ Advanced health check settings

Port

The port the load balancer uses when performing health checks. The default is the traffic port of the load balancer, but you can specify a different port.

☐ Traffic port

☒ Override

8080

1-65535



Select next page

Click on your instance

Register targets

Available instances (1/2)

 *Filter resources by property or value*

	Instance ID		Name
<input type="checkbox"/>	i-02445fb34d9528bdc		JumpHost
<input checked="" type="checkbox"/>	i-04b513102c81868bf		Private01

Click include as pending below

1 selected

Ports for the selected instances

Ports for routing traffic to the selected instances (separate multiple ports with commas):

8080


Include as pending below

Select create group

Create target group

Now create your ALB:

Select Load balancers

 **Load Balancing**

Load Balancers

Select create Load Balancer

Create Load Balancer

Select ALB

Application Load Balancer

HTTP
HTTPS

Create

Choose an Application Load Balancer when you need a flexible feature set for your web applications with HTTP and HTTPS traffic. Operating at the request level, Application Load Balancers provide advanced routing and visibility features targeted at application architectures, including microservices and containers.

[Learn more >](#)

Name the load balancer

Load balancer name

Name must be unique withi

ALB

A maximum of 32 alphanun

Select Internet facing

Scheme [Info](#)

Scheme cannot be changed after the load balancer is created.

☒ **Internet-facing**

An internet-facing load balancer routes requests from clients over the internet to targets. Requires a public subnet. [Learn more](#) 

☐ **Internal**

An internal load balancer routes requests from clients to targets using private IP addresses.

Select IPv4

IP address type [Info](#)

Select the type of IP addresses that your subnets use.

☒ **IPv4**

Recommended for internal load balancers.

☐ **Dualstack**

Includes IPv4 and IPv6 addresses.

Select your VPC

Network mapping [Info](#)

The load balancer routes traffic to targets in the selected subnets, and in accordance with your IP address settings.

VPC [Info](#)

Select the virtual private cloud (VPC) for your targets. Only VPCs with an internet gateway are enabled for selection. The selected VPC cannot be changed. To confirm the VPC for your targets, view your [target groups](#) .

kura-vpc
vpc-0923116d7924ad7d0
IPv4: 192.168.0.0/16



Select two AZ's and two public subnets


☒ **us-east-1a**

Subnet

subnet-033073c2283225dd8

public01 ▼



The subnet for your internet-facing load balancer must have a route to an internet gateway. You can update the subnet's route table in the [VPC Console](#) .

IPv4 settings

Assigned by AWS


☒ **us-east-1b**

Subnet

subnet-0660ccbd79e22236e

public02 ▼



The subnet for your internet-facing load balancer must have a route to an internet gateway. You can update the subnet's route table in the [VPC Console](#) .

Select the security group for the ALB

Security groups [Info](#)

A security group is a set of firewall rules that control the traffic to your load balancer.

Security groups

Select security groups

[Create new security group](#)

Allow_all sg-06e5357bdb155c1db X
VPC: vpc-0923116d7924ad7d0

Select HTTP and your target group

Listeners and routing [Info](#)

A listener is a process that checks for connection requests, using the protocol and port you configure. Traffic received by the listener is then routed per your specifications. You can specify multiple rules and multiple certificates per listener after the load balancer is created.

▼ Listener HTTP:80

Protocol

HTTP ▼

Port

80

1-65535

Default action [Info](#)

Forward to

Jenkins

Target type: Instance

HTTP ▼

[Create target group](#)

Finally select create load balancer

Create load balancer

Takes a few minutes to set up. You will see the status is active and the target group health is healthy in the target group section

Edit only_jumpshot

Custom TCP port 8080

Source 0.0.0.0

Type
Info

Custom TCP

Protocol
Info
Port range
Info

TCP
8080

Source
Info

Custom

0.0.0.0/0

Allow_all security group

Type
Info
Protocol
Info
Port range
Info
Source
Info

All traffic
All
All
Custom

0.0.0.0/0

Go into Load Balancer and copy the DNS name and paste it in url

DNS name

ALB-935408808.us-east-1.elb.amazonaws.com

Create another EC2 inside the same private subnet of the Jenkins master (This will be the agent)

Network ⓘ vpc-0923116d7924ad7d0 | kura-vpc ↕ [Create new VPC](#)

Subnet ⓘ subnet-0119f0ebc4a0b7c31 | private01 | us-east-1a ↕ [Create new subnet](#)
16378 IP Addresses available

Auto-assign Public IP ⓘ Use subnet setting (Disable) ↕

Tag

Name - "Private01 - Child"

Create a new security group

Assign a security group: ☒ Create a **new** security group
☐ Select an **existing** security group

Security group name: jenkins_child

Description: This will allow only children of the jenkins master

Type ⓘ	Protocol ⓘ	Port Range ⓘ	Source ⓘ
SSH ▾	TCP	22	Custom ▾ sg-043377a0f24abb656
Custom TCP F ▾	TCP	8080	Custom ▾ 0.0.0.0/0, ::/0

For SSH Source. Put the security group that has jenkins on it. This is only_jumphost

SSH into JumpHost (Public01),

Once inside, then SSH into private01.

Create a new key and put the RSA information into it

nano linux.pem

chmod 400 linux.pem

Once inside there, SSH once again into private01-child

ssh -i linux.pem ec2-user@Private IPv4 addresses

```

PS C:\Users\robin\.ssh> ssh -i .\rixardo.pem ec2-user@54.82.69.201
Last login: Sat Sep 11 20:07:07 2021 from cpe-24-193-149-222.nyc.res.rr.com

  _| _|_ )
 _| ( _| /  Amazon Linux 2 AMI
__|\_|_|_|

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-192-168-45-222 ~]$ ls
linux.pem
[ec2-user@ip-192-168-45-222 ~]$ ssh -i linux.pem ec2-user@192.168.189.227
Last login: Sat Sep 11 20:07:24 2021 from 192.168.45.222

  _| _|_ )
 _| ( _| /  Amazon Linux 2 AMI
__|\_|_|_|

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-192-168-189-227 ~]$ ls
linux.pem
[ec2-user@ip-192-168-189-227 ~]$ ssh -i linux.pem ec2-user@192.168.185.83
The authenticity of host '192.168.185.83 (192.168.185.83)' can't be established.
ECDSA key fingerprint is SHA256:FRLjRzRdmKbTibyPX/f0/eriCzKU80nSoQ4jSPOxMME.
ECDSA key fingerprint is MD5:58:86:69:35:5b:3d:d7:f5:42:13:7a:f3:ff:2e:5c:3c.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '192.168.185.83' (ECDSA) to the list of known hosts.

  _| _|_ )
 _| ( _| /  Amazon Linux 2 AMI
__|\_|_|_|

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-192-168-185-83 ~]$ █

```

Once on Jenkin's page
 sudo cat /var/lib/jenkins/secrets/initialAdminPassword

Install suggested plugins

Install plugins the Jenkins community finds most useful.

Once that is done, put in information

Configure the Jenkins master to SSH into the agent

Once logged into jenkins, go to Mange Jenkins



Manage Jenkins

Select manage nodes



Manage Nodes and Clouds

Add, remove, control and monitor the various nodes that Jenkins runs jobs on.

Select new node in the left



New Node

Create a name for the node and select permanent agent.

Node name

Test



Permanent Agent

Adds a plain, permanent agent to .
Select this type if no other agent t

OK

Create a name and description

Name

Test

Description

test

Enter 2 for executors


Number of executors

2

Enter {/home/ec2-user/jenkins} for remote root directory

Remote root directory

{/home/ec2-user/jenkins}

 **Are you sure you want to use current working directory. Us**

Create a label - agent-linux

Labels

agent-linux

Select use this node as much as possible

Usage

Use this node as much as possible

Select launch agent via SSH

Launch method


Launch agents via SSH

Enter the private IP address of the agent for Host - 192.168.185.83

Host

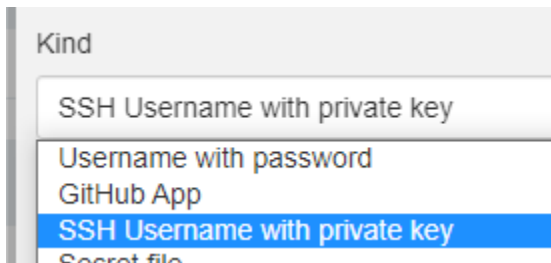
192.168.185.83

Add SSH credentials (username: ec2-user | key: the private key you used to ssh into agent)



For Kind Select

SSH Username with private key



Enter ID - **worker-ssh** and enter description - **ssh into agent**

ID

worker-ssh

Description

ssh into agent

Enter username **ec2-user**

Username

ec2-user

For private key, enter your RSA key information directly into the box.

Private Key

☒ Enter directly

Key

```
JG5WQZKCgYEA14MPCBa1KECqGCqSVLYLIGYZ2n6p0v0FV2KSay0WzHGLyOKpyG8Y
x0IG1PuMntrwi5Re3oJCM1vtA1fHQPRF2K8nqlWs2zpNGwzemHWNjFHMtJxmBIpaQ
4SAgMxNrvhLdyIbakrs2sg7XRMAyRK3F2xXoif5LewPnC8WhlvYt9Qs=
-----END RSA PRIVATE KEY-----
```

No passphrase for the key. Press Add

Select your credentials

Credentials

ec2-user (ssh into agent) ▼

ec2-user (ssh into agent)

Select non verifying verification strategy

Host Key Verification Strategy

Non verifying Verification Strategy

Save and then look at the logs to see if the setup was successful.

Save

Look at the logs to see if your setup was successful

Click on agent





Log

Click on Logs on the left side.

Errors - Java

```
checking Java version in the PATH
bash: java: command not found
Java is not in the PATH nor configured with the javaPath setting, Jenkins will try to guess where is Java, this guess will be removed in the future. :Launch
agents via SSH
[09/11/21 20:34:22] [SSH] Checking java version of {/home/ec2-user/jenkins}/jdk/bin/java
Couldn't figure out the Java version of {/home/ec2-user/jenkins}/jdk/bin/java
bash: {/home/ec2-user/jenkins}/jdk/bin/java: No such file or directory

[09/11/21 20:34:22] [SSH] Checking java version of java
Couldn't figure out the Java version of java
bash: java: command not found

[09/11/21 20:34:22] [SSH] Checking java version of /usr/bin/java
Couldn't figure out the Java version of /usr/bin/java
bash: /usr/bin/java: No such file or directory

[09/11/21 20:34:22] [SSH] Checking java version of /usr/java/default/bin/java
Couldn't figure out the Java version of /usr/java/default/bin/java
bash: /usr/java/default/bin/java: No such file or directory

[09/11/21 20:34:22] [SSH] Checking java version of /usr/java/latest/bin/java
Couldn't figure out the Java version of /usr/java/latest/bin/java
bash: /usr/java/latest/bin/java: No such file or directory

[09/11/21 20:34:22] [SSH] Checking java version of /usr/local/bin/java
Couldn't figure out the Java version of /usr/local/bin/java
bash: /usr/local/bin/java: No such file or directory

[09/11/21 20:34:22] [SSH] Checking java version of /usr/local/java/bin/java
Couldn't figure out the Java version of /usr/local/java/bin/java
bash: /usr/local/java/bin/java: No such file or directory







java.io.IOException: Java not found on hudson.slaves.SlaveComputer@df16141. Install Java 8 or Java 11 on the Agent.
    at hudson.plugins.sshslaves.JavaVersionChecker.resolveJava(JavaVersionChecker.java:84)
    at hudson.plugins.sshslaves.SSHLauncher$1.call(SSHLauncher.java:453)
    at hudson.plugins.sshslaves.SSHLauncher$1.call(SSHLauncher.java:421)
    at java.base/java.util.concurrent.FutureTask.run(FutureTask.java:264)
    at java.base/java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1128)
    at java.base/java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:628)
    at java.base/java.lang.Thread.run(Thread.java:829)
[09/11/21 20:34:22] Launch failed - cleaning up connection
[09/11/21 20:34:22] [SSH] Connection closed.
```

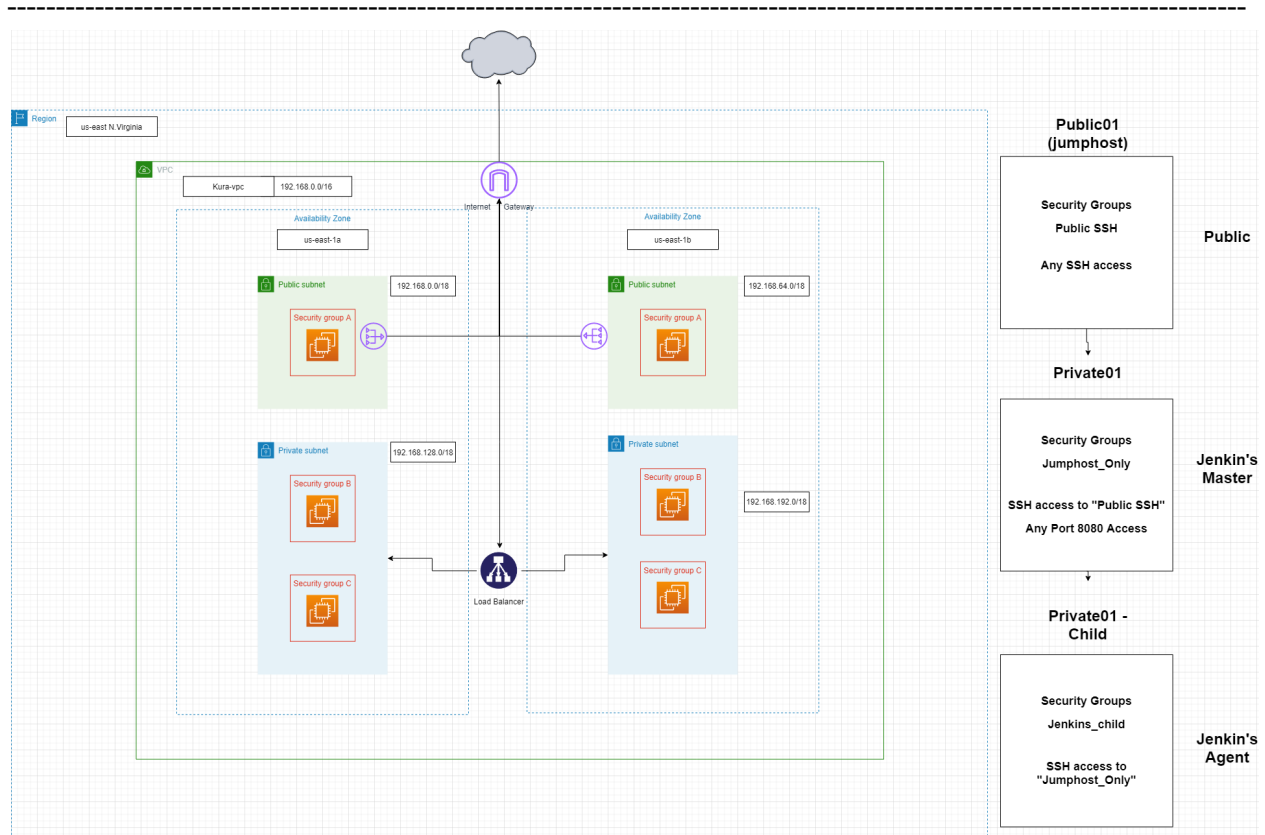
Fix: Install the following in the agent ec2 terminal

```
sudo yum install maven
```

```
sudo yum install git
```

Success:

S	Name ↓	Architecture	Clock Difference	Free Disk Space	Free Swap Space	Free Temp Space	Response Time
	master	Linux (amd64)	In sync	5.90 GB	 0 B	5.90 GB	0ms 
	Test	Linux (amd64)	In sync	5.92 GB	 0 B	5.92 GB	81ms 
Data obtained		1 min 41 sec	1 min 41 sec	1 min 41 sec	1 min 41 sec	1 min 41 sec	1 min 41 sec



So we created a VPC (virtual private cloud) Then we created 4 different subnets with different IP ranges using subnetting. Once we had that we created an internet gateway which basically allows our vpc/devices to communicate outside of the network. In other words it's like a router. Then we created two routing tables for publicRT and privateRT. The routing table is for traffic inside the VPC. For the public subnets, we associated the public routing tables to them and vice versa.

Then we created a public EC2 called jumphost which is the public EC2. This ec2 instance has a public subnet and allows auto-assign public IP.

We then created a private ec2 instance that has the private01 subnet attached. For the security group we assigned the SSH access to the public01 security group. So that basically means that you can only SSH into private01 if you are inside of the public01 (jumphost) instance. Then we created a pem keypair and pasted the information into the file and chmod it. Once inside the private01 instance we pinged and got no response.

We then created a NAT gateway to allow us to restrict inbound but allow us to do outbound for updating. We attached the NAT gateway to the publicRT which basically gave us a response when we pinged

Then we created another EC2 private instance in the same private01 subnet and named it child. For the security group, the SSH's value is the only_jumphost security group. So basically the only_jumphost(private01) can ssh into the child ec2 instance.

User -> Internet/Cloud -> Internet Gateway -> EC2 with attached Nat Gateway -> private01 jenkins

So for this topology there are two AZ zones. They both have a public subnet and a private subnet. Both of them can ssh into the public subnet and the private subnet has a master jenkins program on it. When someone accesses the application load balancer, traffic is either gone to us-east-1a or us-east-1b. This is in case one zone is overloaded and for redundancy. Once inside, the target group redirects traffic from the public subnet to the private subnet that has jenkins installed. It uses port forwarding to direct the traffic.